NovaCode on Terminal Bench 2

Performance report | Recorded Modal evaluations through 14 October 2026 | Code revision 7301cce

49 / 89	55.1%	40
Historical tasks passed	Observed pass rate	Trials without a pass

What the full run measured

The 13 October Terminal Bench 2 run completed all 89 task trials with one attempt per task. Harbor reported 49 passing tasks. The local trace analyzer classified the 40 non-passes as 27 task-verifier failures, 9 execution errors, and 4 graph cutoffs. These categories describe observed outcomes, not root causes.

Outcome	Count	Share of 89	Interpretation
Passed verifier	49	55.1%	Task reward recorded as a pass
Task failure	27	30.3%	Verifier did not accept the solution
Execution error	9	10.1%	Agent or trial failed before a valid pass
Graph cutoff	4	4.5%	Step limit stopped the run

Run profile

Dataset: terminal-bench/terminal-bench-2; model: deepseek-v4.1-flash:cloud; environment: Modal. Recorded usage: 260.4 million input tokens (245.4 million cached; 15.0 million uncached) and 2.09 million output tokens. Token totals are usage counters, not billed cost; the result file has no USD cost.

Trace signal, with a strong caveat

The analyzer found no recognized post-edit check in 20 of 27 verifier failures. It found the same signal in 33 of 49 passes. The detector misses custom scripts and does not judge check relevance, so this pattern cannot be read as a failure predictor or proof that missing checks caused failures.

Changes tested against task outcomes

Small matched replays provide directional evidence, not a new benchmark score.

Matched task replays

Task / treatment	Before	After	What was observed
sparql-university / post-edit reminder	0	1	Matched replay passed; the gate was later revised and left opt-in.
distribution-search / post-edit reminder	1	1	Matched replay remained a pass.
financial-document-processor / numeric + classification guidance	0	1	Intermediate numeric-only run remained 0; final run passed 7/7 verifier tests.

For the financial-document task, the baseline took 3m28s and used 653,817 input / 5,755 output tokens. The final run took 4m56s and used 1,893,768 input / 12,380 output tokens; most final input tokens were cached. It corrected the reported totals and document classification. This is one attempt per variant, so it does not establish a general reliability gain.

What changed in NovaCode

The agent prompt now tells NovaCode to define numeric fields precisely and check ambiguous classifications against the full category definition. An optional post-edit gate prompts once when a recognized check is missing or failed; it is enabled with NOVA_VERIFY_AFTER_EDIT=1 and remains off by default. The analyzer records check attempts, and the job comparator rejects mismatched model, environment, or task profiles.

Practical reading

The 55.1% figure is the historical full-run result, not a measured score for the newly shipped revision. The matched two-task replay and the single financial-document trial are too small to revise that score. A larger repeated, paired panel should measure pass rate, execution failures, graph cutoffs, token cost, and added turns before enabling the reminder by default.

Other exploratory tasks

A chess-best-move replay failed after a board-reading error. A filter-js-from-html replay failed, but a clean-input verifier comparison appears inconsistent with byte-preserving output; its failed XSS tests still prevent any claim of success. Neither was counted as evidence of an improvement.

Sources and reproducibility

Local result folders: nova-evaluation/jobs/tb2/2026-10-13__21-51-25, 2026-10-14__11-28-45, and 2026-10-14__13-16-36 through 13-26-31. Analysis: nova-evaluation/scripts/analyze.py. Matched-profile comparison: nova-evaluation/scripts/compare_jobs.py. Evidence notes: coding-reliability-evidence-2026-10-14.md and coding-reliability-followup-evidence-2026-10-14.md. No live model call was made to prepare this PDF.